

Lab 9

4. Run the commands inside KSQL

CREATE STREAM ConfluentTextLinesStream (line VARCHAR) WITH (KAFKA_TOPIC='ConfluentTextLinesTopic', VALUE_FORMAT='KAFKA');

1

CREATE STREAM ConfluentTextLinesStream (line VARCHAR) WITH (KAFKA_TOPIC='ConfluentTextLinesTopic', VALUE_FORMAT='KAFKA');

• Add query properties

auto.offset.reset

=

Earliest

🗑

+Add another field

Stop

Run query

1

{

2

"@type": "currentStatus",

3

"statementText": "CREATE STREAM CONFLUENTTEXTLINESSTREAM (LINE STRING) WITH (CLEANUP_POLICY='delete',

4

KAFKA_TOPIC='ConfluentTextLinesTopic', KEY_FORMAT='KAFKA', VALUE_FORMAT='KAFKA');",

5

"commandId": "stream/`CONFLUENTTEXTLINESSTREAM`/create",

6

"commandStatus": {

7

"status": "SUCCESS",

8

"message": "Stream created",

9

"queryId": null

10

},

11

"commandSequenceNumber": 6,

12

"warnings": []

13

}

SELECT * FROM ConfluentTextLinesStream EMIT CHANGES;

1

SELECT * FROM ConfluentTextLinesStream EMIT CHANGES;

• Add query properties

auto.offset.reset

=

Earliest

🗑

+Add another field

Running...

Stop

🔄

Data structure

STREAM

Total messages

--

Messages/sec

--

Total message bytes

--



▼ {"LINE":"\\beta\\\"}"}

▼ {"LINE":"\\alpha\\\"}"}

SELECT Ucase(line) FROM ConfluentTextLinesStream EMIT CHANGES;

1

SELECT Ucase(line) FROM ConfluentTextLinesStream EMIT CHANGES;

• Add query properties

auto.offset.reset

=

Earliest

🗑

+Add another field

Running...

Stop

Data structure

STREAM

Total messages

--

Messages/sec

--

Total message bytes

▶ ||

🔍 Filter by keyword



▼ {"KSQL_COL_0":"\BETA\""}

▼ {"KSQL_COL_0":"\ALPHA\""}

CREATE STREAM UppercasedTextLinesStream WITH (KAFKA_TOPIC='UppercasedTextLinesTopic', VALUE_FORMAT='KAFKA') AS SELECT Ucase(line) FROM ConfluentTextLinesStream;

1

CREATE STREAM UppercasedTextLinesStream WITH (KAFKA_TOPIC='UppercasedTextLinesTopic', VALUE_FORMAT='KAFKA') AS SELECT Ucas

• Add query properties

auto.offset.reset

=

Earliest

🗑

+Add another field

Stop

Run query

```
1 {
2   "@type": "currentStatus",
3   "statementText": "CREATE STREAM UPPERCASEDTEXTLINESSTREAM WITH (CLEANUP_POLICY='delete',
4     KAFKA_TOPIC='UppercasedTextLinesTopic', PARTITIONS=10, REPLICAS=3, RETENTION_MS=604800000, VALUE_FORMAT='KAFKA') AS
5     SELECT UCASE(CONFLUENTTEXTLINESSTREAM.LINE) KSQL_COL_0\nFROM CONFLUENTTEXTLINESSTREAM CONFLUENTTEXTLINESSTREAM\nEMIT
6     CHANGES;",
7   "commandId": "stream/'UPPERCASEDTEXTLINESSTREAM'/create",
8   "commandStatus": {
9     "status": "SUCCESS",
10    "message": "Created query with ID CSAS_UPPERCASEDTEXTLINESSTREAM_7",
11    "queryId": "CSAS_UPPERCASEDTEXTLINESSTREAM_7"
12  },
13  "commandSequenceNumber": 8,
14  "warnings": []
15 }
```

5. Check UppercasedTextLinesTopic topic;

SELECT * FROM UppercasedTextLinesStream EMIT CHANGES;

```
1 SELECT * FROM UppercasedTextLinesStream EMIT CHANGES;
```

[Add query properties](#)

auto.offset.reset = Earliest

+Add another field

Running...

Stop

Data structure
STREAM

Total messages
--

Messages/sec
--

Total message bytes
--

Filter by keyword

▼ {"KSQL_COL_0":"\BETA\""}

▼ {"KSQL_COL_0":"\ALPHA\""}

CREATE STREAM readings (sensor VARCHAR KEY, val DOUBLE, location VARCHAR) WITH (
kafka_topic='readings',
partitions=6,
value_format='JSON');

```
1 CREATE STREAM readings (  
2   sensor VARCHAR KEY,  
3   val DOUBLE,  
4   location VARCHAR) WITH (  
5     kafka_topic='readings',  
6     partitions=6,  
7     value_format='JSON');
```

[Add query properties](#)

auto.offset.reset = Earliest

+Add another field

Stop

Run query

```
1 {  
2   "@type": "currentStatus",  
3   "statementText": "CREATE STREAM READINGS (SENSOR STRING KEY, VAL DOUBLE, LOCATION STRING) WITH  
(CLEANUP_POLICY='delete', KAFKA_TOPIC='readings', KEY_FORMAT='KAFKA', PARTITIONS=6, VALUE_FORMAT='JSON');",  
4   "commandId": "stream/'READINGS'/create",  
5   "commandStatus": {  
6     "status": "SUCCESS",  
7     "message": "Stream created",  
8     "queryId": null  
9   },  
10  "commandSequenceNumber": 10,  
11  "warnings": []  
12 }
```

```

INSERT INTO readings (sensor, val, location) VALUES ('sensor-1', 45, 'wheel');
INSERT INTO readings (sensor, val, location) VALUES ('sensor-2', 41, 'motor');
INSERT INTO readings (sensor, val, location) VALUES ('sensor-1', 42, 'wheel');
INSERT INTO readings (sensor, val, location) VALUES ('sensor-3', 42, 'muffler');
INSERT INTO readings (sensor, val, location) VALUES ('sensor-3', 48, 'muffler');

```

```

1  INSERT INTO readings (sensor, val, location) VALUES ('sensor-1', 45, 'wheel');
2  INSERT INTO readings (sensor, val, location) VALUES ('sensor-2', 41, 'motor');
3  INSERT INTO readings (sensor, val, location) VALUES ('sensor-1', 42, 'wheel');
4  INSERT INTO readings (sensor, val, location) VALUES ('sensor-3', 42, 'muffler');
5  INSERT INTO readings (sensor, val, location) VALUES ('sensor-3', 48, 'muffler');

```

[Add query properties](#)

auto.offset.reset = Earliest

[+Add another field](#)
[Stop](#)
[Run query](#)

Data structure

--

Total messages

--

Messages/sec

--

[▶](#)
[⏸](#)

No new messages

The message browser shows messages that have arrived since this page was opened.

4. Run some queries

CREATE STREAM clean AS SELECT sensor, val, UCASE(location) as location FROM readings EMIT CHANGES;

```

1  CREATE STREAM clean AS SELECT sensor, val, UCASE(location) as location FROM readings EMIT CHANGES;

```

[Add query properties](#)

auto.offset.reset = Earliest

[+Add another field](#)
[Stop](#)
[Run query](#)

```

1  {
2    "@type": "currentStatus",
3    "statementText": "CREATE STREAM CLEAN WITH (CLEANUP_POLICY='delete', KAFKA_TOPIC='pksqlc-n7kwdCLEAN', PARTITIONS=6,
4    REPLICAS=3, RETENTION_MS=604800000) AS SELECT\n READINGS.SENSOR SENSOR,\n READINGS.VAL VAL,\n
5    UCASE(READINGS.LOCATION) LOCATION\nFROM READINGS READINGS\nEMIT CHANGES;",
6    "commandId": "stream/'CLEAN'/create",
7    "commandStatus": {
8      "status": "SUCCESS",
9      "message": "Created query with ID CSAS_CLEAN_11",
10     "queryId": "CSAS_CLEAN_11"
11   },
12   "commandSequenceNumber": 12,
13   "warnings": []
14 }

```

SELECT sensor, val, location FROM readings WHERE sensor='sensor-1' EMIT CHANGES;

1

SELECT sensor, val, location FROM readings WHERE sensor='sensor-1' EMIT CHANGES;


● Add query properties

auto.offset.reset = Earliest 

+Add another field

Running...

Stop



Data structure

STREAM

Total messages

--

Messages/sec

--

Total message bytes



Filter by keyword



▼ {"SENSOR":"sensor-1","VAL":42,"LOCATION":"wheel"}

Newest


▼ {"SENSOR":"sensor-1","VAL":45,"LOCATION":"wheel"}

CREATE STREAM high_readings AS SELECT sensor, val, location FROM clean where val > 42 EMIT CHANGES;

1

CREATE STREAM high_readings AS SELECT sensor, val, location FROM clean where val > 42 EMIT CHANGES;

● Add query properties

auto.offset.reset = Earliest 

+Add another field

Stop

Run query

```
1 {
2   "@type": "currentStatus",
3   "statementText": "CREATE STREAM HIGH_READINGS WITH (CLEANUP_POLICY='delete', KAFKA_TOPIC='pksqlc-n7kwdHIGH_READINGS',
4     PARTITIONS=6, REPLICAS=3, RETENTION_MS=604800000) AS SELECT\n  CLEAN.SENSOR SENSOR,\n  CLEAN.VAL VAL,\n  CLEAN.LOCATION LOCATION\nFROM CLEAN CLEAN\nWHERE (CLEAN.VAL > 42)\nEMIT CHANGES;",
5   "commandId": "stream/'HIGH_READINGS'/create",
6   "commandStatus": {
7     "status": "SUCCESS",
8     "message": "Created query with ID CSAS_HIGH_READINGS_13",
9     "queryId": "CSAS_HIGH_READINGS_13"
10  },
11   "commandSequenceNumber": 14,
12   "warnings": []
}
```

5. Produce some events into readings topic

INSERT INTO readings (sensor, val, location) VALUES ('sensor-3', 36, 'motor');

1

INSERT INTO readings (sensor, val, location) VALUES ('sensor-3', 36, 'motor');

● Add query properties

auto.offset.reset

=

Earliest

⌵

🗑

+Add another field

Stop

Run query

Data structure

--

Total messages

--

Messages/sec

--

Total message bytes

--



Filter by keyword



No new messages

The message browser shows messages that have arrived since this page was opened.

6. Run some queries

CREATE STREAM high_pri AS SELECT sensor, val, UCASE(location) as location
FROM readings where val > 42 EMIT CHANGES;

1

CREATE STREAM high_pri AS SELECT sensor, val, UCASE(location) as location

2

FROM readings where val > 42 EMIT CHANGES;

● Add query properties

auto.offset.reset

=

Earliest

⌵

🗑

+Add another field

Stop

Run query

```
1 {
2   "@type": "currentStatus",
3   "statementText": "CREATE STREAM HIGH_PRI WITH (CLEANUP_POLICY='delete', KAFKA_TOPIC='pksqlc-n7kwdHIGH_PRI',
4     PARTITIONS=6, REPLICAS=3, RETENTION_MS=604800000) AS SELECT\n READINGS.SENSOR SENSOR,\n READINGS.VAL VAL,\n UCASE(READINGS.LOCATION) LOCATION\nFROM READINGS READINGS\nWHERE (READINGS.VAL > 42)\nEMIT CHANGES;",
5   "commandId": "stream/'HIGH_PRI'/create",
6   "commandStatus": {
7     "status": "SUCCESS",
8     "message": "Created query with ID CSAS_HIGH_PRI_15",
9     "queryId": "CSAS_HIGH_PRI_15"
10  },
11  "commandSequenceNumber": 16,
12  "warnings": []
}
```

CREATE STREAM by_location AS SELECT * FROM high_pri
PARTITION BY location EMIT CHANGES;

```
1 CREATE STREAM by_location AS SELECT * FROM high_pri
2 PARTITION BY location EMIT CHANGES;
```

● [Add query properties](#)

auto.offset.reset = Earliest 

[+Add another field](#)

Stop

Run query

```
1 {
2   "@type": "currentStatus",
3   "statementText": "CREATE STREAM BY_LOCATION WITH (CLEANUP_POLICY='delete', KAFKA_TOPIC='pksqlc-n7kwdBY_LOCATION',
4     PARTITIONS=6, REPLICAS=3, RETENTION_MS=604800000) AS SELECT *\nFROM HIGH_PRI HIGH_PRI\nPARTITION BY
5     HIGH_PRI.LOCATION\nEMIT CHANGES;",
6   "commandId": "stream/`BY_LOCATION`/create",
7   "commandStatus": {
8     "status": "SUCCESS",
9     "message": "Created query with ID CSAS_BY_LOCATION_17",
10    "queryId": "CSAS_BY_LOCATION_17"
11  },
12  "commandSequenceNumber": 18,
13  "warnings": []
14 }
```

CREATE STREAM by_val AS SELECT * FROM high_pri
PARTITION BY val EMIT CHANGES;

```
1 CREATE STREAM by_val AS SELECT * FROM high_pri
2 PARTITION BY val EMIT CHANGES;
```

● [Add query properties](#)

auto.offset.reset = Earliest 

[+Add another field](#)

Stop

Run query

```
1 {
2   "@type": "currentStatus",
3   "statementText": "CREATE STREAM BY_VAL WITH (CLEANUP_POLICY='delete', KAFKA_TOPIC='pksqlc-n7kwdBY_VAL', PARTITIONS=6,
4     REPLICAS=3, RETENTION_MS=604800000) AS SELECT *\nFROM HIGH_PRI HIGH_PRI\nPARTITION BY HIGH_PRI.VAL\nEMIT CHANGES;",
5   "commandId": "stream/`BY_VAL`/create",
6   "commandStatus": {
7     "status": "SUCCESS",
8     "message": "Created query with ID CSAS_BY_VAL_19",
9     "queryId": "CSAS_BY_VAL_19"
10  },
11  "commandSequenceNumber": 20,
12  "warnings": []
13 }
```

```
INSERT INTO readings (sensor, val, location) VALUES ('sensor-1', 45, 'wheel');
INSERT INTO readings (sensor, val, location) VALUES ('sensor-2', 41, 'motor');
INSERT INTO readings (sensor, val, location) VALUES ('sensor-1', 42, 'wheel');
INSERT INTO readings (sensor, val, location) VALUES ('sensor-3', 42, 'muffler');
INSERT INTO readings (sensor, val, location) VALUES ('sensor-3', 48, 'muffler');
```

1

INSERT INTO readings (sensor, val, location) VALUES ('sensor-1', 45, 'wheel');

2

INSERT INTO readings (sensor, val, location) VALUES ('sensor-2', 41, 'motor');

3

INSERT INTO readings (sensor, val, location) VALUES ('sensor-1', 42, 'wheel');

4

INSERT INTO readings (sensor, val, location) VALUES ('sensor-3', 42, 'muffler');

5

INSERT INTO readings (sensor, val, location) VALUES ('sensor-3', 48, 'muffler');

● Add query properties

auto.offset.reset

=

Earliest

⌵

🗑

⚡ Add another field

Stop

Run query

Data structure

--

Total messages



Filter by keyword



```
CREATE TABLE avg_readings AS
  SELECT sensor, AVG(val) as avg
  FROM readings
  GROUP BY sensor
  EMIT CHANGES;
```

1

CREATE TABLE avg_readings AS

2

SELECT sensor, AVG(val) as avg

3

FROM readings

4

GROUP BY sensor

5

EMIT CHANGES;

● Add query properties

auto.offset.reset

=

Earliest

⌵

🗑

⚡ Add another field

Stop

Run query

```
1  {
2    "@type": "currentStatus",
3    "statementText": "CREATE TABLE AVG_READINGS WITH (CLEANUP_POLICY='compact', KAFKA_TOPIC='pksqlc-n7kwdAVG_READINGS',
PARTITIONS=6, REPLICAS=3, RETENTION_MS=604800000) AS SELECT\n READINGS.SENSOR SENSOR,\n AVG(READINGS.VAL) AVG\nFROM
READINGS READINGS\nGROUP BY READINGS.SENSOR\nEMIT CHANGES;",
4    "commandId": "table/AVG_READINGS/create",
5    "commandStatus": {
6      "status": "SUCCESS",
7      "message": "Created query with ID CTAS_AVG_READINGS_21",
8      "queryId": "CTAS_AVG_READINGS_21"
9    },
10   "commandSequenceNumber": 22,
11   "warnings": []
12 }
```



```
CREATE TABLE part_avg AS
  SELECT location, AVG(val) as avg
  FROM readings
  GROUP BY location
  EMIT CHANGES;
```

```
1 CREATE TABLE part_avg AS
2   SELECT location, AVG(val) as avg
3   FROM readings
4   GROUP BY location
5   EMIT CHANGES;
```

● [Add query properties](#)

auto.offset.reset = Earliest 

[+Add another field](#)

Stop

Run query

```
1 {
2   "@type": "currentStatus",
3   "statementText": "CREATE TABLE PART_AVG WITH (CLEANUP_POLICY='compact', KAFKA_TOPIC='pksqlc-n7kwdPART_AVG',
4     PARTITIONS=6, REPLICAS=3, RETENTION_MS=604800000) AS SELECT\n  READINGS.LOCATION LOCATION,\n  AVG(READINGS.VAL)
5     AVG\nFROM READINGS READINGS\nGROUP BY READINGS.LOCATION\nEMIT CHANGES;",
6   "commandId": "table/'PART_AVG'/create",
7   "commandStatus": {
8     "status": "SUCCESS",
9     "message": "Created query with ID CTAS_PART_AVG_23",
10    "queryId": "CTAS_PART_AVG_23"
11  },
12  "commandSequenceNumber": 24,
13  "warnings": []
14 }
```